

DMPP: main concepts

Bruno Rocha Pereira

June 14, 2016

Chapter 1

Introduction

- **Deutsch fallacies (you don't need to know the list by heart but to understand what each fallacy implies).**

In a perfect world:

1. The network is reliable!
2. The network is secure!
3. The network is homogeneous!
→ All devices have the same configuration, ports and OS
4. The topology does not change!
5. Latency is zero!
6. Bandwidth is infinite!
7. Transport cost is zero!
8. There is one administrator

- **Distribution vs concurrency, distribution vs. mobility.**

Distributed systems are concurrent programs with no resource shared. Parallel programs can be used with only one core while still having concurrency problem.

Mobile as in mobile phone. Distributed as in using a system with devices on the same network sending each other messages.

- **Distribution transparency and its fundamental issues (again, you don't need to know the list of problems by heart but to understand what each issue implies and why distribution transparency is said to be a myth).**

- Latency: There can be a delay between sending a message and the moment when the other device gets it.
- Difference in Memory Access : Every method invocation implies sending information over the network to retrieve the reference.
- Partial Failure: An element of the system can fail
- Concurrency: automatic parallelisation cannot be made automatic because of partial failure and latency

- **Language vs. middleware vs. reflective approach.**

- Language : concepts create a language, hiding the details
- Middleware : library, makes use of another language's functionalities.
- Reflective Approach : use of Aspects or reflection to combine both

- **Essential complexity vs. accidental complexity.**

- Essential complexity: complexity bound to the problem. Cannot be avoided
- Accidental complexity: bound to the development of the solution. Can be avoided by the dev.

Chapter 2

Prototype-based Programming

- Why are prototypes relevant for distributed programming?
- Symbiosis: understand the concept, how symbiosis works between AmbientTalk and Java.
- Object creation exnihilo vs. cloning, cloning and instantiation, delegation and cloning.
- Scoping: lexical vs object scope, methods vs. closures.
- Forwarding vs. delegation.
- The uniform access principle (UAP), UAP and closures, firstclass methods.
- Classifying objects with type tags.
- The concept of traits, objects as traits.

Chapter 3

Event-loop Programming

- Threads: understand the model, how to achieve inter and intraobject synchronization, understand why and how deadlocks happen.
- Actors: the basic principles of the model, and difference with the thread model.
- Active Objects: understand the model, and its limitations for distribution.
- Event Loops: understand the model, the concept of far references, and the three concurrency control properties that the model can enforce.
- Actors in AmbientTalk, asynchronous message sending and parameter passing rules.
- Asynchronous message sending and return values:
- The concept of customer objects, and its limitations.
- Futures: the concept, blocking vs. nonblocking futures, future pipelining, how to achieve conditional synchronization with futures.

Chapter 4

Event-based Distributed Programming

- Understand the different design issues to provide a distributed object model, and its relation to the different types of networks.
- Know the discriminating properties of mobile networks, and understand the requirements for a distributed object model for MANETs.
- Far references and partial failures: how far references deal with intermittent failures.
- Ways of obtaining a first far reference, service discovery vs. service lookup.
- Distributed object scoping and isolates.
- Far references and permanent failures, leased object references.
- Understand the difference between pessimistic and optimistic distributed object model.

Chapter 5

Meta-level Engineering

- Metaprogramming: firstclass messages, quasiquoting and splicing, first-class abstract grammar
- Reflective programming: base vs meta level, meta vs. reflective program.
- Terminology on reflection (reification, introspection, intercession), structural vs. behavioural reflection.
- The concept of meta objects, and understand the problems with popular meta-level architectures
- Mirror-based reflection: understand the model and the three properties that the model can enforce.
- Mirrors in AmbientTalk: explicit vs. implicit mirrors on objects.
- Mirages: the concept.
- Mirrors on Actors: understand which operations actor mirrors reify in comparison to object mirrors.
- MOP: You are not expected to know by heart the message invocation protocol but the overall idea and understand the relevant parts that need to be altered in order to implement language constructs like futures and leased references.

Chapter 6

Coordination using Tuple Spaces

- Terminology: coordination, datadriven vs. control driven coordination.
- Tuple Spaces: the basic interaction mechanism, understand the Linda model and how tuple matching works.
- Know what decoupling in time and space, and synchronization decoupling is, and being able to reason about the forms of decoupling that tuple spaces and other communication paradigms seen at the course exhibit.
- Federation of mobile tuple spaces: understand the basic interaction model of tuple spaces in a mobile setting, and being able to explain the two variations (LIME & TOTA), know what are the differences and similarities of the two variations.
- TOTAM: understand the goal for this hybrid approach, and key features, and how the model supports memory management in face of failures.

6.1 Decoupling

6.1.1 Space decoupling

The interacting parties do not need to know each other. The publishers publish events through an event service, and the subscribers get these events indirectly through the event service. Publishers and subscribers do not hold references to each other neither do they know how many of them are participating

6.1.2 Time decoupling

The interacting parties do not need to be actively participating in the same interaction at the same time. Events occurring when a subscriber is offline are delivered when he gets back online.

6.2 Tuple space

6.2.1 Primitives

out an actor outputs a tuple into the tuplespace

peek an actor reads matching tuples

in an actor grabs a tuple into itself

6.2.2 LIME (Linda In a Mobile Environment)

*Put more intelligence in the tuplespace **operations** (*in, out, peek, ...*)*

This is a federated tuple space model (make everyone agree). Hard to guarantee consistency: atomicity only for **in** operations. Implies tradeoffs for decoupling in space and time.

6.2.3 TOTA (Tuples Over The Air)

*Put more intelligences in the **tuples** themselves*

A replication-based tuple space model. Fully decoupled in space and time, but atomicity cannot be guaranteed.

- Tuples are injected into the network, and hop between spaces.
- Each tuple carries a propagation and a maintenance rule; ie **Tuple** = (**Content**, **propRule**, **maintRule**)
- **in/peek** only affects the local tuplespace and direct neighbors

Supporting Tuples and safe state

T1 in a space S1 is a **supporting tuple** of T2 in a space S2 if:

- They share the same key
- There is 1 hop between S1 and S2
- $T1.hop = n$ and $T2.hop = n+1$

A tuple is in **safe state** if it has at least one supporting tuple, or it is a source tuple ($hop = 0$) A tuple **deletes itself** if it is no longer in a safe state

6.2.4 TOTAm (Tuples Over The Ambient)

Mixed approach of LIME and TOTA. Atomicity for **in** operations. **read** decoupled in time. Scoped propagation protocol. Reacts on incoming tuples.

Private and public tuples

The classical out operation acts on the local tuplespace. **inject** into the ambient. Each tuple carries a propagation protocol and a sign. To remove tuples, inject an **antituple**: a tuple with the same content and propagation rules, but opposite sign.

Chapter 7

Peer-to-peer systems

- Typical characteristics from P2P systems.
- The concept of an overlay network.
- First generation of P2P systems: motivation and limitations.
- Second generation of P2P systems: understand the basics, flooding, TTL propagation
- Understand the motivation for P2P third generation and which guarantees they provide w.r.t. previous generations, distributed hash tables
- Chord: understand how keys are distributed, the lookup algorithm and the different cost models of lookup, how join and leave of nodes affect the network, and the role of periodic stabilization
- CAP theorem: understand consistency, availability, partition tolerance means and their tradeoff. Be able to reason which guarantee Chord and other distributed algorithms or applications cannot provide.
- Distributed storage: know why P2P systems scale well specially for immutable objects, being able to reason about different replication strategies.
- Beernet: the goal of the approach and basic architecture. Understand the differences with Chord with respect to management of peer failures, and why it is relevant in a mobile setting.

Fundamentals

- Understand why external synchronization with physical clocks is not employed for coordinating distributed processes.
- Lamport's clocks: the algorithm, know which properties Lamport's clocks exhibit, understand what it means that Lamport clocks are not strongly consistent.
- Partial vs. total ordering of events.
- Vector clocks: Understand the difference with Lamport's clocks, which properties they exhibit, and how you can compare vector clocks. Understand the limitations and assumptions of vector clocks.
- Be able to reason about application domains where lamport and vector clocks can be applied.
- Consensus: two Generals' problem, know what the consensus problem is.
- Know what it means the safety and liveness properties, and being able to reason about those properties for distributed algorithms.
- Know the FLP Theorem and understand its implications.
- Paxos:
 - understand the goal, the roles of nodes and the purpose of each phase.
 - Be able to explain why the algorithm cannot reach agreement sometimes, and how the safety properties are guaranteed.
 - Be able to reason about application domains where paxos is useful.

Chapter 8

Distributed Programming on the Web

- Know the basic interaction model with a client-server architecture communicating by means of HTTP request.
- Understand why programming on the web 1.0 is said to be programming with strings.
- Know what a RIA is, and the role of JavaScript in the development of RIAs.
- AJAX: understand the idea that scripts can asynchronously communicate with the server, and why the interactions run asynchronously.
- Thread-based vs. event-based servers.
- Know the two ways how to achieve communication amongst clients (Publish/ Subscribe on top of HTTP and websockets) and reason about their advantages and disadvantages.
- Mobile cross-platform solutions:
 - The idea of employing web-based technologies (i.e. JavaScript) for mobile computing, and which advantages brings to mobile software development.
 - Understand the basics of two big family of cross-platform solutions (hybrid vs. interpreted) and be able to reason about their advantages and disadvantages.

Chapter 9

Reactive Programming

- Understand the differences between classic sequential and event-driven software.
- Know the advantages of event-driven client and servers for RIAs.
- Know what the issue of “Inversion of Control” (a.k.a. the callback hell effect) is, and its implications (w.r.t shared state)
- Reactive programming: the basic model of evaluation, and key abstractions (event sources vs. behaviours), lifting, glitches and the most common dataflow evaluation strategy to reevaluate code without causing glitches.
- Flapjax:
 - understand how pages can be made reactive
 - you do not need to know all the operations on event stream combinators, but you should be able to reason about a piece of code with respect to which parts are reactive, the dependency graph, and how reactive code communicates with the DOM and server).
 - Understand what it means that there are no callbacks in Flapjax in contrast to JavaScript.